

## 1 Qualify the evidence

### What?

Inspect datasets, labels and time records.

### Why?

Confirm that each source supports the question.

### How?

Record versions, class counts and dependencies.

## 2 Prepare the evaluation

### What?

Build current-flow, role and history features.

### Why?

Compare methods on the same later records.

### How?

Use forward folds and matched class budgets.

## 3 Fit the comparison models

### What?

Fit flow experts and evidence selectors.

### Why?

Test how information choices change decisions.

### How?

Use fixed LightGBM settings and three seeds.

## 4 Record error destinations

### What?

Measure correct stage, wrong stage and benign.

### Why?

Expose missed warnings behind overall scores.

### How?

Report stage recall, warning recall and false alerts.

## 5 Check uncertainty and scope

### What?

Reanalyse paired predictions and source support.

### Why?

Show variation, repeated evidence and limits.

### How?

Use capture resampling and seed/capture omissions.

## 6 Deliver the praxis evidence

### What?

Assemble findings, figures, tables and code.

### Why?

Make the evaluation procedure inspectable.

### How?

Verify public arithmetic and bind artifact hashes.